28 inputs → 64 tanh → 64 tanh → 1 linear force output

6,081 parameters | normalized force × 650 N, clipped to [0, 650] N

$$L_{ordinary} = L_{current} \qquad L_{CL} = L_{current} + \beta L_{memory}$$

Each L = weighted normalized force MSE + 0.1 corrective-response penalty
Replay: 1,024 target-task TRAIN examples; at most 512 per optimizer update